Zhuoran Lu

Updated October 15, 2026

Email: zhuoranlu33@gmail.com **Website**: https://zhuoranlu.github.io/ **Phone**: +1 (217)-419-2369

RESEARCH INTERESTS
Human-AI Interaction, Human-centered AI, Human-AI Collaboration, Responsible AI, Agentic AI, Applied Machine Learning, Computational Social Science

EDUCATION

Purdue University, West Lafayette, IN Aug 2019 - May 2025
Ph.D. in Computer Science
Advisor: Dr. Ming Yin.
Committee: Dr. Ninghui Li, Dr. Tianyi Zhang, Dr. Sooyeon Jeong.

University of Illinois at Urbana-Champaign, Champaign, IL Jan 2017 - May 2019
B.S. in Computer Science and Statistics; Minor in Mathematics
Graduated with Highest Distinction
Advisor: Dr. Roxana Girju and Dr. Richard Sowers

APPOINTMENTS

Accenture Center for Advanced AI May 2025 - Current
Research Scientist Senior Manager, Center of Advanced AI, *Mountain View, CA*
Worked with: Dr. Wei Wei & Dr. Michelle Zhou

Autodesk Research May 2024 - Aug 2024
Research Intern, From Intent to Experience Group, *Toronto, ON*
Worked with: Dr. Yi Wang & Dr. Qian Zhou

Microsoft Research May 2023 - Sep 2023
Research Intern, Office of Applied Research, *Redmond, WA*
Worked with: Dr. Mengting Wan & Dr. Longqi Yang

Honda Research Institute Jan 2023 - Apr 2023
Research Intern, Human-centered Intelligence Group, *San Jose, CA*
Worked with: Dr. Shashank Mehrotra

National Center of Supercomputing Applications Jan 2018 - May 2018
Software Engineer Intern, Advanced Reactors and Fuel Cycles, *Urbana, IL*
Worked with: Dr. Kathryn D. Huff

PUBLICATIONS * indicates equal contribution, † indicates senior author

Peer-reviewed Conference Papers

[C.24] Yibo Meng, Bingyi Liu, Lionel Z. Wang, and **Zhuoran Lu**†. Digital Privacy in End-of-Life Care: A Qualitative Study of Terminally Ill Patients, Family Caregivers, and Clinicians. In *Proc. of the 48th IEEE Symposium on Security and Privacy (IEEE S&P)*, Montreal, May 2027.

[C.23] Jie Gao, Kaiser Sun, Jen-tse Huang, Katherine Van Koevering, Sijie Ji, Heyuan Huang, Weiyan Shi, **Zhuoran Lu**†, Ziang Xiao, Daniel Khashabi, and Mark Dredze. How to Interpret Agent Behavior. In *Proc. of the 40th Conference on Neural Information Processing Systems (NeurIPS)*, Sydney, December 2026.

[C.22] Lingyi Long, Xian Li, Shuai Ma, Yi Wang, and **Zhuoran Lu**†. FANVERSE: Exploring and Extending Branching Story Universes Across Canon and Fan Fiction. In *the Creative AI Track of the 40th Conference on Neural Information Processing Systems (NeurIPS)*, Sydney, December 2026.

[C.21] **Zhuoran Lu**, Yangyang Yu, Zhuoyan Li, Yibo Meng, Nan Jiang, Chengxi Zang, Jie Gao, and Ziang Xiao. Using Grounded Theory for Agent Behavior Analysis at Scale. In *Findings of the 2026 Conference on Empirical Methods in Natural Language Processing (EMNLP Findings)*, Budapest, October 2026.

[C.20] **Zhuoran Lu**, Weilong Wang, Yangyang Yu, Xinru Wang, Zhuoyan Li, Zhiwei Liu, and Sophia Ananiadou. One AI Signal, Many Human Judgments: A Bayesian Cascade Analysis of AI-based Credibility Indicators in Online Information Spread. In *Proc. of the 2026 ACM Conference on Human-AI Complementarity and Alignment (HCOMP)*, Washington DC, September 2026. **[Best Paper Honorable Mention]**

[C.19] Zhuoyan Li, Ze Ma, Shishuang He, **Zhuoran Lu**, Ziang Xiao, and Ming Yin. Adaptive Selection of Effective AI Assistance in AI-assisted Decision Making Using Multi-Armed Bandits. In *Proc. of the 2026 ACM Conference on Human-AI Complementarity and Alignment (HCOMP)*, Washington DC, September 2026. **[Best Paper Honorable Mention]**

[C.18] **Zhuoran Lu***, Gionnieve Lim*, and Ming Yin. Large Language Model (LLM)-driven Adversarial Social Influences in Online Information Spread: Risks and Interventions. In *Proc. of the 44th ACM Conference on Human Factors in Computing Systems (CHI)*, Barcelona, May 2026.

[C.17] Zhuoyan Li, Aditya Bansal, Jinzhao Li, Shishuang He, **Zhuoran Lu**, Mutian Zhang, Qin Liu, Yiwei Yang, Swati Jain, Ming Yin, and Yunyao Li. Human-LLM Collaborative Feature Engineering for Tabular Data. In *Proc. of the 14th International Conference on Learning Representations (ICLR)*, Rio de Janeiro, April 2026.

[C.16] **Zhuoran Lu**, Patrick Li, Weilong Wang, Ming Yin. Understanding the Effects of AI-based Credibility Indicators When People Are Influenced By Both Peers and Experts. In *Proc. of the 43rd ACM Conference on Human Factors in Computing Systems (CHI)*, Yokohama, May 2025.

[C.15] **Zhuoran Lu**, Qian Zhou, Yi Wang. WhatELSE: Shaping Narrative Space at Configurable Level of Abstraction for AI-bridged Interactive Storytelling. In *Proc. of the 43rd ACM Conference on Human Factors in Computing Systems (CHI)*, Yokohama, May 2025.

[C.14] Zhuoyan Li, Hangxiao Zhu, **Zhuoran Lu**, Ziang Xiao, Ming Yin. From Text to Trust: Empowering AI-assisted Decision Making with Adaptive LLM-powered Analysis. In *Proc. of the 43rd ACM Conference on Human Factors in Computing Systems (CHI)*, Yokohama, May 2025.

[C.13] **Zhuoran Lu**, Dakuo Wang, Ming Yin. Does More Advice Help? The Effects of Second Opinions in AI-Assisted Decision Making. In *Proc. of the 27th ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW)*, San José, November 2024.

[C.12] **Zhuoran Lu**, Syed Hasan Amin Mahmood, Zhuoyan Li, Ming Yin. Mix and Match: Characterizing Heterogeneous Human Behavior in AI-assisted Decision Making. In *Proc. of the 12th AAAI Conference on Human Computation and Crowdsourcing (HCOMP)*, Pittsburgh, October 2024.

[C.11] **Zhuoran Lu***, Syed Hasan Amin Mahmood*, Ming Yin. Designing Behavior-Aware AI to Optimize Human-AI Team Decision Making. In *Proc. of the 32nd International Joint Conference on Artificial Intelligence (IJCAI)*, Jeju, August 2024.

[C.10] Chun-wei Chiang, **Zhuoran Lu**, Zhuoyan Li, Ming Yin. Enhancing AI-Assisted Group Decision Making through LLM-Powered Devil's Advocate. In *Proc. of the 29th ACM Conference on Intelligent User Interfaces (IUI)*, Greenville, March 2024.

[C.9] Zhuoyan Li, **Zhuoran Lu**, Ming Yin. Decoding AI's Nudge: A Unified Framework to Predict Human Behavior in AI-assisted Decision Making. In *Proc. of the 38th Conference on Artificial Intelligence (AAAI)*, Vancouver, February 2024.

[C.8] Zhuoyan Li, Hangxiao Zhu, **Zhuoran Lu**, Ming Yin. Synthetic Data Generation with Large Language Models for Text Classification: Potential and Limitations. In *Proc. of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Singapore, December 2023.

[C.7] **Zhuoran Lu***, Zhuoyan Li*, Chun-Wei Chiang, Ming Yin. Strategic Adversarial Attacks in AI-assisted Decision Making to Reduce Human Trust and Reliance. In *Proc. of the 32nd International Joint Conference on Artificial Intelligence (IJCAI)*, Macao, August 2023.

[C.6] Chun-Wei Chiang, **Zhuoran Lu**, Zhuoyan Li, Ming Yin. Are Two Heads Better Than One in AI-Assisted Decision Making? Comparing the Behavior and Performance of Groups and Individuals in Human-AI Collaborative Recidivism Risk Assessment. In *Proc. of the 41st ACM Conference on Human Factors in Computing Systems (CHI)*, Hamburg, May 2023.

[C.5] **Zhuoran Lu***, Zhuoyan Li*, Ming Yin. Modeling Human Trust and Reliance in AI-assisted Decision Making: A Markovian Approach. In *Proc. of the 37th Conference on Artificial Intelligence (AAAI)*, Washington D.C., February 2023. **[Oral Presentation]**

[C.4] **Zhuoran Lu**, Patrick Li, Weilong Wang, Ming Yin. The Effects of AI-based Credibility Indicators on the Detection and Spread of Misinformation under Social Influence. In *Proc. of the 25th ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW)*, Online, November 2022. **[Best Paper Award]**

[C.3] **Zhuoran Lu***, Zhuoyan Li*, Ming Yin. Towards Better Detection of Biased Language with Scarce, Noisy, and Biased Annotations. In *Proc. of the 5th AAAI-ACM Conference on AI, Ethics, and Society (AIES)*, Oxford, August 2022.

[C.2] **Zhuoran Lu***, Xinru Wang*, Ming Yin. Will You Accept the AI Recommendation? Predicting Human Behavior in AI-Assisted Decision Making. In *Proc. of the ACM Web Conference (WWW)*, Online, April 2022.

[C.1] **Zhuoran Lu**, Ming Yin. Human Reliance on Machine Learning Models When Performance Feedback is Limited: Heuristics and Risks. In *Proc. of the 39th ACM Conference on Human Factors in Computing Systems (CHI)*, Online, May 2021.

Peer-reviewed Journal Articles

[J.4] Xinru Wang, **Zhuoran Lu**, Shuai Ma, Ming Yin, Eunyee Koh, and Mustafa Doga Dogan. XAgen: Designing Explainability Support for Failure Detection and Correction in Multi-Agent Workflows. In *Design and Artificial Intelligence (DAAI)*, 2026.

[J.3] Yibo Meng, Xinru Wang, Zhuoyan Li, Chengxi Zang, **Zhuoran Lu**†, and Zhicong Lu. When Ubiquitous Health Data Stops Helping: Patients' Experience of Breakdowns in the Data-Interpretation-Action Loop at the End of Life. In *Proc. of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)*, 2026.

[J.2] Yibo Meng, Jinyuan Zhan, Yao Lyu, **Zhuoran Lu**†, and Chengxi Zang. "Can Computing Systems Eventually Disappear?": A Long-Term Tracking Study on the Use of Diabetes Computing Systems by Patients with Diabetes. In *Proc. of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies (IMWUT)*, 2026.

[J.1] Shan Gao, **Zhuoran Lu**, Hao Luan, Ming Yin, Lei Wang. AI Pilot in the Cockpit: An Investigation of Public Acceptance. In *International Journal of Human-Computer Interaction (IJHCI)*, 2024.

Preprints

[P.6] Yulin Chen, Yang Zhan, **Zhuoran Lu**†, and Qiao Jin. How LLM Assistants Come to Involve People Beyond Their User. *Preprint*

[P.5] Hasan Amin, **Zhuoran Lu**, and Ming Yin. Likelihood Is Not Authorship: Can We Reliably Measure Human Contribution in AI-Assisted Writing? *Preprint*

[P.4] Ismail Hossain, Sai Puppala, **Zhuoran Lu**†, Sajedul Talukder, and Nan Jiang. Benchmarking Security Risk Detection and Verification in Open Agentic Skill Ecosystems. *Preprint*

[P.3] Zhiwei Liu, Yangyang Yu, Yupeng Cao, Yuechen Jiang, Haohang Li, **Zhuoran Lu**†, Yuyan Wang, Yixiang Zheng, Xiaorui Guo, Calvin Yixiang Cheng, and Sophia Ananiadou. AutoRedTrader: Autonomous Red Teaming of Trading Agents through Synthetic Misinformation Injection. *Preprint*

[P.2] Yibo Meng, Ruiqi Chen, **Zhuoran Lu**†, Shuai Ma, and Chengxi Zang. Tracing Generative AI in Digital Art: A Longitudinal Study of Chinese Painters' Attitudes, Practices, and Identity Negotiation. *Preprint*

[P.1] **Zhuoran Lu**, Sheshera Mysore, Tara Safavi, Jennifer Neville, Longqi Yang, Mengting Wan. Augmenting Workplace Social Media Composition via Personalized LLM-powered Writing Assistance. *Preprint*

Lightly Reviewed Workshop Papers

[W.8] Shixiang You, Zeyuan Wang, Yuechen Jiang, Zhiwei Liu, Jie Gao, Sophia Ananiadou, and **Zhuoran Lu**†. GreenAudit-Bench: Evidence-Grounded Auditing of Compositional Greenwashing in Corporate Sustainability Disclosures. In *the EMNLP Workshop on Financial Technology and Natural Language Processing (FinNLP)*, Budapest, October 2026.

[W.7] Yibo Meng, Bingyi Liu, Yao Lyu, Jinyuan Zhan, and **Zhuoran Lu**†. When Technologies Disappear, So Does the User: A Longitudinal Study of Wearable Health Computing. In *Companion of the 2026 ACM International Joint Conference on Pervasive and Ubiquitous Computing (UbiComp Companion)*, Shanghai, October 2026.

[W.6] Yibo Meng, Bingyi Liu, Xinru Wang, Zhuoyan Li, and **Zhuoran Lu**†. Persistent Signals, Absent Meaning: How Terminal-Stage Patients Navigate Health Data Exposure in Everyday Ubiquitous Computing Environments. In *Companion of the 2026 ACM International Joint Conference on Pervasive and Ubiquitous Computing (UbiComp Companion)*, Shanghai, October 2026.

[W.5] Yuxuan Du, Shuai Ma, and **Zhuoran Lu**†. What Personalized Agents Make Visible: A Diagnostic Framework for AI-Mediated Judgment. In *Companion of the 2026 ACM International Joint Conference on Pervasive and Ubiquitous Computing (UbiComp Companion)*, Shanghai, October 2026.

[W.4] **Zhuoran Lu***, Gionnieve Lim*, Ming Yin. Understanding the Effects of Large Language Model-driven Adversarial Social Influence in Information Spread. In *Proc. of the Late-Breaking Work (LBW) at the 43rd ACM Conference on Human Factors in Computing Systems (CHI)*, Yokohama, May 2025.

[W.3] Sheshera Mysore, **Zhuoran Lu**, Mengting Wan, Longqi Yang, Steve Menezes, Tina Baghaee, Emmanuel Barajas Gonzalez, Jennifer Neville, Tara Safavi. PEARL: Personalizing Large Language Model Writing Assistants with Generation-Calibrated Retrievers. In *the EMNLP Customizable NLP Workshop*, Miami, November 2024.

[W.2] **Zhuoran Lu***, Syed Hasan Amin Mahmood*, Ming Yin. Give Weight to Human Reactions: Optimizing Complementary AI in Practical Human-AI Teams. In *the ICML AI and HCI Workshop*, Hawaii, July 2023.

[W.1] **Zhuoran Lu**, Dakuo Wang, Ming Yin. Does More Advice Help? The Effects of Second Opinions from Peers in AI-Assisted Decision Making. In *the CHI Workshop on Trust and Reliance in AI-Assisted Tasks (TRAIT)*, Hamburg, May 2023.

PATENT

[P.T.1] Yi Wang, **Zhuoran Lu**, Qian Zhou. Narrative Space Generation for AI-Bridged Interactive Storytelling. US Patent Application 2026/0115599 A1

HONORS & SCHOLARSHIPS

Best Paper Honorable Mention, ACM HCOMP [C.20]	2026
Best Paper Honorable Mention, ACM HCOMP [C.19]	2026
Bilsland Dissertation Fellowship, Purdue University	2024
Special Recognition for Outstanding Reviews, ACM CHI	2023
Best Paper Award, ACM CSCW [C.4]	2022
Frederick N. Andrews Fellowship, Purdue University	2019
Best Project of the Year, Illinois Geometry Lab, UIUC	2018

TEACHING EXPERIENCE

Graduate Teaching Assistant, Purdue University

CS 373: Data Mining & Machine Learning	Fall 2021, Spring 2022
CS 180: Object-oriented Programming	Fall 2019

Guest Lecturer, Purdue University

CS 373: Data Mining & Machine Learning	Spring 2022

· Topic: Trustworthy Machine Learning - A Human-centric View

CS 590: Human-AI Interaction — Fall 2022, Spring 2024

· Topic: Human Reliance on AI - Internal and External Signals

INVITED TALKS

Accounting Individual Cognition & Social Influence for Better Human-AI Collaboration

Data and Systems Engineering Seminar, The University of Hong Kong	Apr 2026
HCI Seminar, Johns Hopkins University	Mar 2025
Responsible AI Seminar, Carnegie Mellon University	Mar 2025
Info Technology & Management Seminar, Illinois Institute of Technology	Mar 2025
Computer Science Seminar, Illinois Institute of Technology	Mar 2025
Industrial & Sys Engineering Colloquium, University of Wisconsin-Madison	Feb 2025
HCI Colloquium, Indiana University at Indianapolis	Feb 2025
Computer Science Seminar, University of Houston	Feb 2025
Human-centered AI Design Seminar, Northeastern University	Feb 2025

Human Reliance on Machine Learning - Internal and External Signals

Information Management Seminar, Nanjing University	Dec 2023
Information Management Young Scholar Forum, Nanjing University	Dec 2022

Personalized LLM-empowered Writing Assistant for Workplace Social Media

Microsoft Research — Sep 2023

Experiments and User Studies in Information Credibility Research

Information Science Methodology Seminar, Wuhan University — Aug 2023

SERVICE

Organizer

Interpret Agent Behavior (IAB) Workshop @ NeurIPS: 2026

Associate Chair

ACM Conference on Human Factors in Computing Systems (CHI): 2027

ACM CHI Late-Breaking Work: 2024, 2025, 2026

Program Committee Member

ACM Conference on Fairness, Accountability, and Transparency (FAccT): 2023, 2024, 2025

ACM International Conference on Conversational User Interfaces (CUI): 2024

ACM International Joint Conference on Pervasive and Ubiquitous Computing (Ubicomp): 2023

Conference Reviewer

Conference on Neural Information Processing Systems (NeurIPS): 2025, 2026

Empirical Methods in Natural Language Processing (EMNLP): 2025, 2026

ACM Conference on Human Factors in Computing Systems (CHI): 2023, 2024, 2025, 2026

AAAI Conference on Web and Social Media (ICWSM): 2023, 2024

ACM Conference on Computer-Supported Cooperative Work and Social Computing (CSCW): 2023, 2024, 2025
ACM Conference on Designing Interactive Systems (DIS): 2022, 2023, 2024, 2026
ACM Conference on Intelligent User Interfaces (IUI): 2023, 2024, 2025
ACM Interaction Design and Children (IDC) Conference: 2023, 2024
Meeting of the Cognitive Science Society (CogSci): 2023, 2024
Computer-Human Interaction of Australia (OzCHI): 2024
ACM Symposium on User Interface Software and Technology (UIST): 2023
Symposium on Computer-Human Interaction in Play (CHI PLAY): 2023
ACM International Conference on Conversational User Interfaces (CUI): 2023
ACM International Conference on Mobile Human-computer Interaction (MobileHCI): 2022
ACM Symposium on Spatial User Interaction (SUI): 2022
International Conference on Human-Agent Interaction (HAI): 2022
Nordic Human-Computer Interaction Conference (NordiCHI): 2022
ACM International Conference on Multimodal Interaction (ICMI): 2022
ACM Interactive Surfaces and Spaces Conference (ISS): 2023
ACM International Joint Conference on Pervasive and Ubiquitous Computing (Ubicomp): 2022

Journal Reviewer
ACM Transactions on Knowledge Discovery from Data (TKDD)
ACM Transactions on Interactive Intelligent Systems (TiiS)
International Journal of Human-Computer Interaction (IJHCI)
Information Processing and Management (IPM)
Human Factors and Ergonomics in Manufacturing & Service Industries (HFE)
Computational Brain & Behavior

MENTORING EXPERIENCE

Yulin Chen, M.S. student at Carnegie Mellon University
Yibo Meng, M.S. student at Cornell University
Yuxuan Du, M.S. student at Southeast University
Xian Li, M.Phil. student at HKUST (Guangzhou)
Xiaokang Lei, M.Phil. student at HKUST (Guangzhou)
Lingyi Long, B.S. student at Shanghai Jiao Tong University
Hangxiao Zhu, M.S. student at Washington University in St. Louis
Tony Xu, B.S. student at Purdue University
Taylor Gong, B.S. student at University of Illinois at Urbana-Champaign